

The Drift Complexity Index: Matching Detector Resolution to Workload Drift

Arun Reungsilpolkarn

arun.r@bu.ac.th

School of Information Technology and Innovation, Bangkok University
Pathum Thani, Thailand

Abstract

Workload drift—the motion of a workload’s feature vector—has a measurable *complexity*: the effective number of independent modes the motion occupies. We introduce the *Drift Complexity Index* (DCI), a participation ratio of the drift-motion covariance, computable in $O(D^2)$ without eigendecomposition. We prove that DCI bounds detector sufficiency. The best 1-D detector retains the dominant drift mode’s share of the multi-dimensional detector’s power, and DCI brackets that share: a cheap detector provably suffices below one frontier and provably cannot above a second. Between the two, the spectrum decides. Empirically, ordinary drift is *near* rank-one, so one dimension suffices; only heterogeneous drift needs the full kernel. The regime is a property of the drift, not the kernel: kernel-free features reproduce it, and a real 100k-query trace exhibits it. The application: index advisors are gated behind an always-on drift detector. DCI decides, before any detector runs, whether the cheap one is safe. Live on PostgreSQL and MongoDB, a fixed 1-D detector structurally misses heterogeneous drift — recall 0.60, stale indexes at 10.9× query latency. The DCI-routed selector matches the always-multi-dimensional detector decision for decision, at 38% of its monitoring cost over the offline suite. Where full resolution is required, DCI announces it in advance. DCI is a monitoring-resolution primitive for self-driving data systems. Advisor gating is its first application.

Keywords

workload drift, drift detection, index advisor, workload similarity, complexity index

1 Introduction

Modern databases ship index advisors that recommend a physical design for a workload, but an advisor is expensive to run: it searches a large space of hypothetical indexes, often estimating costs for thousands of candidates. Running one on every window of a long-lived workload is wasteful — most of the time the workload has not changed enough to warrant new advice. The standard remedy is to *gate* the advisor: invoke it only when the workload *drifts*. Gating in turn needs a drift detector that — unlike the advisor it gates — runs on *every* window. It is the always-on monitoring layer.

How elaborate must that layer be? A rich, multi-dimensional workload-similarity kernel is safe — it sees drift along any axis — but it is expensive to evaluate every window. A single cheap similarity dimension is the opposite: nearly free, but blind to drift that does not move that axis. Choosing between them blindly is unsatisfying. This paper asks whether one can tell, cheaply and

before committing to a detector, how much detection machinery a workload’s drift requires.

Our answer is that workload drift has a measurable *complexity*. Drift is the motion of a workload’s feature vector; the number of independent modes that motion occupies is a workload property, estimated at runtime with the *Drift Complexity Index* (DCI) — a participation ratio of the drift-motion covariance, computable in $O(D^2)$ with no eigendecomposition. DCI is revealing. Ordinary workload drift is *near* rank-one: a feature launch or a traffic change moves the similarity axes together, so a single dimension already captures it. Genuinely complex drift requires *heterogeneous* change — several kinds of drift at once — and only then is the full kernel needed. Because DCI separates these regimes *before* any detector runs, it routes the monitoring layer to the cheapest sufficient detector. We treat DCI not as a one-off gating heuristic but as a general primitive—a cheap, engine-portable measure of *how* a workload drifts—whose first application is deciding a detector’s resolution. The layer it builds is deliberately *advisor- and engine-agnostic*: it reads only query text and arrival timestamps and treats the advisor as a black box, so the same construction runs unchanged in front of any advisor, on any engine whose workload yields a per-window feature vector—validated on a relational and a document store.

This paper makes four contributions. **(C1)** The Drift Complexity Index: a cheap, bounded, engine-portable *primitive* for workload-drift complexity, with the eigenvalue bounds that tie it to detector sufficiency (§4). **(C2)** A *regime map*: across three SQL workloads and a document store (MongoDB), a single similarity axis suffices below a DCI frontier near 1.5 and the full kernel is needed above it (§6); the same regime is reproduced by kernel-free feature representations and appears on a real workload trace (§6.8). **(C3)** A DCI-routed selector with a derived, verified sufficiency tolerance; it retains 99% of the full kernel’s detection-AUC gain at 38% of its monitoring cost (§5, §6). **(C4)** A practitioner rule, validated end-to-end: profile a workload’s DCI and deploy the matching detector; a compact deployment on live PostgreSQL and MongoDB (§6.9) confirms the routed advisor closes the loop in both regimes — on mixed drift it matches the always-multi-D detector’s decisions where a fixed 1-D detector structurally cannot. The full economic value of gating — throughput, plan quality — composes with but is separable from ours, and is left to future work (§7).

2 Related Work

Index advisors — the gated component. The advisor DCI gates itself a deep and active line of work. Cost-driven index selection and the optimiser-coupled “what-if” interface [5, 6] matured into productised advisors that jointly choose indexes, views, and partitions [1], and into online, learning-based tuners: bandit formulations with regret-based safety guarantees [26, 27] and advisors for dynamic, heterogeneous workloads [34]. Recent work pushes on scale [4, 33], online uncertainty [35], learned search over hybrid designs [24], and LLM-guided configuration [15]. An in-depth

study consolidates the index-advisor design space [38], and the self-driving-DBMS programme frames the broader goal [25]. This machinery is increasingly *engine-portable*: cross-DBMS tuners span PostgreSQL and MySQL [31, 32], and cost-based physical design reaches NoSQL document and column stores [20]. All of this work decides *which* physical design to build, or *how* to search for it. DCI sits one layer above and is advisor-agnostic: it decides, cheaply and on every window, *how much detection machinery* a workload’s drift warrants before the (expensive) advisor is consulted at all. Even a drift- or uncertainty-aware advisor [34, 35] must first be reached; DCI governs the detector that decides when to reach it.

Drift and change-point detection. The always-on layer DCI routes is a drift detector. Concept-drift detection — DDM [8], ADWIN [3], and the surveys that organise a still-growing field [2, 9, 18] — and the sequential change-point theory beneath it — CUSUM [23] and window-limited GLR schemes with delay/false-alarm optimality [13, 14] — answer *whether* and *when* drift occurred, and a large recent literature continues to sharpen their accuracy and adaptivity [2]. DCI answers a prior, orthogonal question: *how complex* the drift is, and hence which detector to run at all. Lu et al. [18] name “drift understanding” as a task distinct from detection; DCI turns that qualitative notion into a concrete, $O(D^2)$ statistic that routes the monitoring layer. That the cost of the always-on detector is itself a first-class concern is increasingly recognised [22], and drift is becoming a first-class *benchmarking* dimension: DriftBench generates data- and workload-drift for benchmark inputs [17] — the generation-side complement of DCI, which prices a given drift’s complexity.

Hierarchical hypothesis testing. The nearest prior work is hierarchical hypothesis testing for drift [36, 37]: a cheap, always-on Layer-I test backed by an expensive Layer-II test invoked *after* a Layer-I alarm. That is *escalate-on-event* — the expensive machinery runs in response to a detected event. DCI is *select-on-complexity*: it picks the detector *before* any window is scored, from a stationary property of the workload’s drift geometry. The two are orthogonal and composable — DCI sets the resolution of the always-on layer, and an escalation step could still confirm its alarms.

Effective dimensionality. DCI’s core, the participation ratio, originates in condensed-matter physics as a count of effective modes [16] and appears in computational neuroscience as the effective dimensionality of neural covariance [10]. To our knowledge this is its first use to characterise *workload* drift and to route a database monitoring layer. It is likewise distinct from workload *characterization and forecasting*, which models what a workload *will be* [12, 19]: DCI measures a structural property of *how* a workload drifts.

3 Background: The HSM Workload-Similarity Kernel

The drift signal that DCI operates on comes from an HSM-style workload-similarity feature [29, 30], which is *prior work*; we define the instantiation we use in full here only so the paper is self-contained. Our contributions — the Drift Complexity Index, its bounds (§4), and the DCI-routed selector (§5) — take the kernel as nothing more than a black-box source of the per-window feature vector, and apply to any other source of one. The specific five axes are one instantiation in the HSM tradition [29, 30], not a contribution of this paper; §6.8 confirms DCI’s regime is

reproduced by a generic similarity that uses none of them. (That several axes are needed to *span* the drift a kernel can meet across workloads [29, 30] is compatible with any one block’s drift being near rank-one: spanning concerns the directions drift *can* take, DCI how many a given block’s drift *does* occupy.)

HSM is a *pre-execution* measure: it reads a workload’s query text and the arrival timestamps of its queries, and never executes a query, builds an index, or inspects an execution plan. A sliding *window* of queries is aggregated into a per-template frequency vector, a query count, the sets of tables and columns referenced, and a one-second-resolution arrival series. From a pair of windows the kernel computes five pairwise similarity scores, each in $[0, 1]$ (1 = identical; Table 1), and a weighted composite. A window’s feature vector $f = (S_R, S_V, S_T, S_A, S_P)$ is thus its similarity to a reference window.

Definition 1 (Window and feature vector). A *window* W_t is the query text and arrival timestamps of a sliding block of queries. Fixing a reference window W_0 , the kernel maps W_t to $f(t) = (S_R, S_V, S_T, S_A, S_P) \in [0, 1]^5$, the five pairwise similarity scores of W_t against W_0 (Table 1), each equal to 1 at identity and decreasing with dissimilarity. *Workload drift* is the motion of $f(t)$ across t ; every construct in this paper is a function of that trajectory.

For the rest of the paper, then, each window is simply a bounded, correlated feature vector $f \in [0, 1]^5$, and that vector is all the detectors ever see. The 1-D detector (§5) tracks drift on a *single* axis of f ; the multi-dimensional detector tracks all five jointly, so the routing question is simply how many of the five a workload’s drift actually moves. Nothing downstream depends on the kernel *itself*: DCI is a function of the covariance of f ’s motion alone, so a plain per-template query-frequency vector or a learned query/plan embedding is an equally valid input; only the numerical frontier (§6) is feature-set-specific. This lets the *same* construction read a document store in §6.2: MongoDB aggregation-pipeline templates and their collections and fields take the place of SQL templates and tables and columns, while the five-axis vector is computed identically.

The five axes span distinct facets of a window pair — template ranking and direction (S_R, S_T), query volume (S_V), the tables and columns referenced (S_A), and the arrival-time shape (S_P); Table 1 gives their closed forms, and the composite weights are the fixed defaults of the HSM line [29, 30].

Two kernel properties matter for what follows. First, the five axes are *correlated*: S_R and S_T are two readings of the same per-template frequency vector — a rank correlation and a cosine angle — and the measured Gram off-diagonals on a real workload run from 0.68 to 0.91. A detector that fuses the axes must therefore account for this covariance rather than treat them as independent (§5). Second, the composite HSM is a fixed weighted *average*; it therefore attenuates drift confined to a single axis, so we use it as a gating score, not a detector (§5).

4 The Drift Complexity Index

Drift is the *motion* of the feature vector f . Over a block of windows we form the $D \times D$ covariance C of the drift-motion deviations ($D = 5$): a symmetric positive-semidefinite matrix with eigenvalues $\lambda_1 \geq \dots \geq \lambda_D \geq 0$ and normalised spectrum $p_i = \lambda_i / \sum_j \lambda_j$. The *Drift Complexity Index* is

$$\text{DCI} = \frac{\text{tr}(C)^2}{\|C\|_F^2} = \frac{1}{\sum_i p_i^2}. \quad (1)$$

Table 1: The five similarity axes and the composite of the HSM-style feature used in our experiments [29, 30]; each score is in $[0, 1]$.

Axis	Formula	Measures
S_R	$1 - \arccos(\rho_s)/\pi$	template rank corr.
S_V	$\min(Q_a, Q_b)/\max(Q_a, Q_b)$	query-count ratio
S_T	$1 - \frac{\pi}{2} \arccos(\hat{v}_a^\top \hat{v}_b)$	template cosine
S_A	$\frac{1}{2}J(T_a, T_b) + \frac{1}{2}J(C_a, C_b)$	table/column Jaccard
S_P	$\sum_{\beta=1}^4 \lambda_\beta sc_\beta$	arrival-pattern shape
HSM	$.25S_R + .20S_V + .20S_T + .20S_A + .15S_P$	composite

This is the *participation ratio* of the drift spectrum — equivalently the exponentiated order-2 Rényi entropy [28] or the Hill number of order two [11]: the effective number of independent modes the drift occupies.

DCI has the three properties a runtime index needs. It is *bounded* — $\text{DCI} \in [1, D]$, equal to 1 when the drift is rank-one and to D when its energy is spread evenly — so a threshold on it is interpretable. It is *cheap*: the closed form in (1) is a trace and a Frobenius norm, $O(D^2)$ with no eigendecomposition — the spectrum itself would cost $O(D^3)$. Measured on the pinned platform (identical values to 10^{-13}), the closed form costs 1.4 vs. the spectrum’s 4.9 μs at this paper’s $D=5$, and 59 vs. 6157 μs at $D=500$ — a $105\times$ ratio that matters exactly where richer per-window features (raw frequency vectors, learned query embeddings; §6.8) would take DCI. computed once per block to set the route, not per window like a detector, so this cost is amortized off the monitoring hot path. And it is *dimensionless*, a ratio of like-unit quantities, hence comparable across workloads and engines without calibration.

Two exact bounds on the dominant mode’s share $p_1 = \max_i p_i$ ground the regime claim. Recall $p_i \geq 0$ and $\sum_i p_i = 1$, and write $S = \sum_i p_i^2 = \text{DCI}^{-1}$.

Proposition 1 (low complexity forces a dominant mode). $p_1 \geq \text{DCI}^{-1}$, with equality iff every non-zero p_i equals p_1 .

PROOF. Each $p_i \leq p_1$ and $p_i \geq 0$, so $p_i^2 = p_i \cdot p_i \leq p_i p_1$; summing over i , $S = \sum_i p_i^2 \leq p_1 \sum_i p_i = p_1$, i.e. $p_1 \geq \text{DCI}^{-1}$. Equality forces $p_i \in \{0, p_1\}$ for every i . $\square$

Proposition 2 (high complexity forbids one). $p_1 \leq \text{DCI}^{-1/2}$, with equality iff the drift is rank-one.

PROOF. Every term of S is non-negative, so $S = \sum_i p_i^2 = p_1^2 + \sum_{i \geq 2} p_i^2 \geq p_1^2$; hence $p_1 \leq \sqrt{S} = \text{DCI}^{-1/2}$. Equality forces $p_i = 0$ for all $i \geq 2$. $\square$

The two bounds bracket the dominant mode. As $\text{DCI} \rightarrow 1$, Proposition 1 drives $p_1 \rightarrow 1$: a single mode carries the drift and a one-dimensional detector captures it. As DCI grows, Proposition 2 drives p_1 down: no axis can carry the drift, so a one-dimensional detector must lose information. Both bounds are tight — attained at the flat-support and rank-one spectra. DCI therefore predicts, before any detector runs, which regime a workload’s drift is in.

Proposition 3 (complexity counts concurrent drift causes). *If, over a block, the drift is driven by k mutually uncorrelated causes — $d(t) = \sum_{c=1}^k g_c(t) v_c$ with fixed directions v_c and uncorrelated zero-mean amplitudes g_c — then $C = \sum_c \text{Var}(g_c) v_c v_c^\top$ has rank at most k and $\text{DCI} \leq k$.*

PROOF. C is a sum of k rank-one PSD terms, so $\text{rank}(C) \leq k$ and at most k of the p_i are non-zero. For a distribution on r

non-zero atoms, $\sum_i p_i^2 \geq 1/r$ by Cauchy–Schwarz, so $\text{DCI} = 1/\sum_i p_i^2 \leq r \leq k$. $\square$

Proposition 3 explains the regime map mechanistically: a single-cause change—a template shift or a volume change—gives $k=1$ and rank-one drift, whereas drift on two independent fronts gives $k=2$ and $\text{DCI} \leq 2$. Measured values sit just under this ceiling (§6: single-cause cells at $\text{DCI} \approx 1$, mixed at ≈ 2), the small excess being finite-sample noise. The two regimes are thus not an artifact of a particular generator but reflect how many independent causes drive the drift—a property of the workload, not the kernel.

4.1 A worked example

The bounds turn a single scalar into a routing decision. Under *template* drift on TPC-H the drift covariance is almost rank-one ($\text{DCI} = 1.02$); Proposition 1 guarantees $p_1 \geq 0.98$, so one axis carries $\geq 98\%$ of the drift energy and the cheap 1-D detector suffices—indeed three axes are within tolerance of the full kernel. Under *mixed* drift ($\text{DCI} = 1.97$), Proposition 2 caps $p_1 \leq 0.71$, so no single axis tracks the drift and the multi-dimensional detector is required. At the frontier $\tau = 1.5$, Propositions 1–2 bracket $p_1 \in [0.67, 0.82]$ —neither dominant (as at low DCI) nor forbidden (as at high DCI), the transition zone the empirical frontier occupies.

4.2 From spectrum to detector: a sufficiency guarantee

The bounds constrain the drift spectrum; we now connect that spectrum to what a detector can see. The intuition is one line: a linear detector’s power is its signal-to-noise ratio along the direction it watches, so, averaged over random drifts, the best fixed direction captures exactly the dominant eigenvalue’s share of the drift energy — p_1 . Formally, model the steady state as $d \sim (0, \Sigma)$ and a drift as a mean shift s , random over a block with $s \sim (0, C)$; a linear detector along a unit direction a has deflection (squared SNR) $\rho(a; s) = (a^\top s)^2 / (a^\top \Sigma a)$, and the multi-dimensional detector realises the full non-centrality $\rho_{\text{full}} = s^\top \Sigma^{-1} s$.

Theorem 1 (dimensional deflection ratio). *Under white steady noise ($\Sigma \propto I$), averaged over drifts $s \sim (0, C)$, the best fixed 1-D detector’s mean deflection relative to the full detector’s mean non-centrality equals the dominant-mode share,*

$$R = \frac{\max_a \mathbb{E}_s [\rho(a; s)]}{\mathbb{E}_s [\rho_{\text{full}}]} = \frac{\lambda_1}{\sum_i \lambda_i} = p_1,$$

hence $\text{DCI}^{-1} \leq R \leq \text{DCI}^{-1/2}$ by Propositions 1–2.

PROOF. Under $\Sigma \propto I$, $\mathbb{E}_s [\rho(a; s)] = (a^\top C a) / (\sigma^2 \|a\|^2)$, maximised at the top eigenvector of C with value λ_1 / σ^2 ; and $\mathbb{E}_s [\rho_{\text{full}}] = \text{tr}(C) / \sigma^2$. The ratio is $\lambda_1 / \text{tr}(C) = p_1$; the bounds are Propositions 1–2 on the spectrum of C . $\square$

The identity is neutral—one direction captures a fraction *exactly* p_1 of the drift energy—and the routing regime follows as a corollary and its converse.

Corollary 1 (sufficiency domain). For $\delta \in (0, 1)$, on $\{\text{DCI} \leq 1/(1 - \delta)\}$ the best 1-D detector retains $R = p_1 \geq \text{DCI}^{-1} \geq 1 - \delta$ of the full non-centrality.

Proposition 4 (tightness at integer complexity). For every integer $m \in \{2, \dots, D\}$ there is a drift covariance with $\text{DCI} = m$ for which the best 1-D detector retains only $R = 1/m$; the sufficiency guarantee of Corollary 1 is tight.

PROOF. Take $C = \text{diag}(\frac{1}{m}, \dots, \frac{1}{m}, 0, \dots, 0)$ with m equal non-zero eigenvalues: $p_i = 1/m$, $\text{DCI} = m$, $p_1 = 1/m$, so $R = 1/m$ by Theorem 1. (The restriction to integer m is forced: by Proposition 1, $R = 1/\text{DCI}$ requires a flat spectrum, whose support size is an integer.) $\square$

Corollary 2 (a sufficiency trichotomy). Fix $\delta \in (0, 1)$. If $\text{DCI} \leq (1 - \delta)^{-1}$, the best 1-D detector is $(1 - \delta)$ -sufficient (Corollary 1). If $\text{DCI} > (1 - \delta)^{-2}$, no single direction is: $R = p_1 \leq \text{DCI}^{-1/2} < 1 - \delta$. In between, the bracket $[\text{DCI}^{-1}, \text{DCI}^{-1/2}]$ straddles $1 - \delta$, so the bounds alone do not decide: sufficiency is spectrum-dependent, with both behaviours realised.

PROOF. The first zone is Corollary 1; the second follows from Proposition 2, since $\text{DCI} > (1 - \delta)^{-2} \iff \text{DCI}^{-1/2} < 1 - \delta$. For the middle zone, both behaviours occur at $\delta = \frac{1}{3}$ (zone (1.5, 2.25]): the flat spectrum with $m = 2$ has $\text{DCI} = 2$ and retains only $R = \frac{1}{2} < \frac{2}{3}$, while a spiked spectrum — $p_1 = 0.67$ with four equal tails of 0.0825 — has $\text{DCI} = 2.10$ yet retains $R = p_1 = 0.67 > \frac{2}{3}$. $\square$

The pieces compose into a floor for the routed system as a whole.

Corollary 3 (routed-system floor). Under the router at threshold τ , every window-block is scored by a detector whose mean deflection is at least $\min(1, \tau^{-1})$ of the full non-centrality: blocks routed to the multi-dimensional detector realise it entirely, and blocks routed to the 1-D detector have $\text{DCI} < \tau$, hence $R = p_1 \geq \text{DCI}^{-1} > \tau^{-1}$ by Proposition 1 and Theorem 1.

At the deployed $\tau = 1.5$, the routed monitoring layer never operates below two-thirds of the full detector's mean deflection — on *any* block, including misrouted ones — while paying the cheap tier's price wherever the guarantee permits it.

Theorem 1 turns the empirical frontier into a guarantee for the deployed axis-restricted detector: across the nine cells the measured 1-D-over-full deflection ratio tracks p_1 and the detection-AUC gap widens exactly as p_1 falls (single-cause cells $p_1 \geq 0.83$, gap ≈ 0.00 ; mixed cells $p_1 \approx 0.62$, gap 0.15–0.20). Two honest boundaries. Detection AUC is a *saturating* function of deflection, so the AUC frontier sits *above* the energy domain (a cell can carry $p_1 = 0.82$ of the energy and still reach AUC 1). And the white-noise model matches the deployed detector, which does not whiten; the whitened refinement — the participation ratio of $\Sigma^{-1/2}C\Sigma^{-1/2}$, strictly more powerful where the steady noise is correlated — is measured in §6.7 and prices as the full kernel.

5 Mechanism: Detectors and the DCI Router

Two detectors. Let $d(t) = f(t) - f_{\text{steady}}$ be a window's deviation from the steady-state mean feature. The *1-D detector* scores drift on the single most informative similarity axis, $-d_j(t)$, with the axis j chosen per workload. The *multi-dimensional detector* is the




Figure 1: The DCI-routed monitoring layer. A complexity probe — the Drift Complexity Index, measured on each window-block — routes the always-on detector: low-complexity drift to the cheap 1-D detector, genuinely complex drift to the multi-dimensional detector. Whichever detector is chosen runs on *every* window; only its drift verdict invokes the expensive index advisor.

Mahalanobis distance $d(t)^T \Sigma^{-1} d(t)$ under the steady-window noise covariance Σ ; since the five axes are correlated (§3), this whitening is the statistically optimal linear fusion — under a Gaussian steady-state model it is the Neyman–Pearson statistic [21]. Both detectors read the workload's *position* relative to steady state, not its motion, so each is single-edged on a transient drift dip — a motion detector would fire twice, once on the dip and once on the recovery.

Why not the composite. The scalar HSM is a fixed weighted average: a drift confined to one axis is attenuated by that axis's weight, while the noise on HSM aggregates all five. The composite is therefore a gating score, not a detector; §6 shows its detection AUC collapsing on a workload where a single axis is perfect.

The router. A runtime selector measures DCI on a sliding window-block and routes (Figure 1): $\text{DCI} < \tau$ selects the 1-D detector, $\text{DCI} \geq \tau$ the multi-dimensional detector, with $\tau \approx 1.5$ and a hysteresis fallback to the multi-dimensional detector on any ambiguity; Corollary 3 floors the routed system at τ^{-1} of the full deflection. The router is a *single scalar*: we find (§6) that adding a drift signal-strength coordinate to DCI does not improve routing.

The sufficiency tolerance. Calling the 1-D detector “sufficient” needs a tolerance: how large a multi-D-over-1-D AUC gap is still within noise? We derive it rather than set it by hand. DeLong's nonparametric estimator [7] gives the variance of the (correlated) AUC gap; pooled across seeds and evaluated at a fixed reference sample size N_{ref} , it yields a per-axis tolerance δ_d computed from each workload's own measured variance and stable as the experiment's seed count grows. An axis is sufficient iff its mean gap does not exceed δ_d .

6 Evaluation

We ask four questions. **RQ1:** is workload-drift complexity cheaply and reliably measurable? **RQ2:** does it predict which detector




Figure 2: Detection AUC vs DCI for all 450 runs (50 seeds $\times$ 3 configurations $\times$ 3 workloads). Marker shape and colour denote the drift configuration; open markers are the multi-D detector, filled markers the 1-D detector — no information is colour-alone. At low DCI the 1-D detector ties the multi-dimensional detector; at high DCI (*mixed*) the 1-D detector falls while the multi-dimensional detector holds. Dashed line: the routing threshold $\tau \approx 1.5$.

a workload needs, and where is the frontier? **RQ3:** does a DCI-routed selector preserve detection quality at materially lower monitoring cost? **RQ4:** does the routed selector drive a live index advisor end-to-end?

6.1 Setup

Three SQL workloads exercise the kernel: TPC-H (18 of its 22 query templates — those in the vendored generator’s deterministic parameter pool; Q15, Q16, Q19, and Q22 are absent from it) and the Join Order Benchmark (JOB, 33 template families) are analytic; pgbench (the TPC-B-like five-statement script) is OLTP. All three are presented to the HSM kernel as SQL text. Drift trajectories come from a generator anchored on the phase profiles of a real MongoDB workload trace, under three configurations: *template* drift moves the template-structure axes, *volume* drift moves S_V , and *mixed* drift rotates the two so the pooled drift spans both axis groups. An arrival-only configuration is deliberately absent: the kernel’s arrival axis S_P does not respond to arrival-timing changes while templates are held fixed—a measured property of this feature set, reported rather than hidden—so arrival drift cannot be isolated as a single-cause cell (§7). Each (workload, configuration) cell is run over 50 seeds; seeds are derived deterministically, so a re-run on a second machine reproduced every DCI and AUC to the printed digit. Code and data reproducing every number, figure, and table are available at <https://github.com/stevezhai139/dci-reproducibility>. over five timing repeats, measured on an Apple M4 MacBook Pro (24 GB unified memory, macOS 26) single-threaded under CPython 3.13.3 with NumPy, SciPy, pandas, and the fastdtw and pywavelets libraries. The kernel’s wavelet and DTW path depends on the last two libraries; only the absolute costs, not DCI or the regime map, depend on this platform.

6.2 The regime map (RQ1, RQ2)

Figure 2 plots detection AUC against DCI for all 450 runs. The structure is clean. Ordinary drift is *low-complexity*: template and

Table 2: Per-cell regime map: the DCI and the 1-D-sufficient axes — those within the derived tolerance δ_d of the multi-dimensional detector. Every *mixed* cell has none; every *template/volume* cell has at least one. The first three workloads are relational (SQL); the MongoDB block (document store; identical drift schedule, 50 seeds per cell) reproduces the same structure — for it we list the verified dominant sufficient axis.

Workload	Drift	DCI	1-D-sufficient axes
TPC-H	template	1.02	S_R, S_T, S_A
TPC-H	volume	1.41	S_V, S_P
TPC-H	mixed	1.97	— (multi-D needed)
JOB	template	1.11	S_A
JOB	volume	1.13	S_V, S_P
JOB	mixed	1.92	— (multi-D needed)
pgbench	template	1.03	S_R, S_T
pgbench	volume	1.01	S_V, S_P
pgbench	mixed	1.75	— (multi-D needed)
MongoDB	template	1.04	S_R
MongoDB	volume	1.01	S_V
MongoDB	mixed	1.80	— (multi-D needed)

volume cells sit at $\text{DCI} \in [1.0, 1.4]$, and there the 1-D detector ties the multi-dimensional detector. Heterogeneous (*mixed*) drift is *high-complexity*: $\text{DCI} \approx 1.8$ –2.0, and there the best single axis falls to AUC 0.75–0.82 while the multi-dimensional detector holds 0.91–1.00.

Table 2 gives the per-cell verdict under the derived tolerance δ_d . Every *mixed* cell has no 1-D-sufficient axis — the multi-dimensional detector is needed — and every template/volume cell has at least one. The two regimes are separated by an empty DCI band — the open interval (1.41, 1.75) between the largest low-complexity cell mean and the smallest high-complexity cell mean — so a single threshold $\tau \approx 1.5$ routes them; across the 50 seeds per cell this separation holds at 95% confidence—every low-complexity cell’s DCI 95% upper bound is ≤ 1.48 and every high-complexity cell’s lower bound ≥ 1.68 . The OLTP workload pgbench corroborates the analytic ones: its drift takes the same shape — low-DCI for template and volume, high-DCI for mixed. The bands sit where the trichotomy of Corollary 2 puts them: with the deployed threshold $\tau = 1.5 = (1 - \delta)^{-1}$, i.e. $\delta = \frac{1}{3}$, every low-complexity cell lies in the guaranteed-sufficiency zone and every mixed cell in the indeterminate zone (1.5, 2.25], where the bounds alone do not decide and the experiment does — against a single axis, as the near-flat measured spectra ($p_1 \approx 0.63$ –0.68, Table 3) predict.

Cross-data-model detection: the map transfers to MongoDB. The construction is not tied to SQL. Applying the identical drift schedules to a MongoDB workload of 24 aggregation-pipeline templates (the kernel reads each query’s pre-execution feature vector, so no database is executed) reproduces the same structure (Table 2, MongoDB block): single-axis drift stays low-complexity ($\text{DCI} = 1.04, 1.01$; one axis and the full kernel both at $\text{AUC} \geq 0.998$), while mixed drift is high-complexity ($\text{DCI} = 1.80$; best single axis AUC 0.82 vs. multi-D 1.00)—the same collapse and recovery, in the same DCI bands, routed by the same τ . This is *cross-data-model* evidence at the detection layer; heterogeneous workload types (semi-structured, vector) remain future work.

A real workload sits in the regime. On a real SkyServer query log (99,969 queries, natural drift only) the same $O(D^2)$ probe reads a per-block DCI averaging 1.85 over 100 blocks and ranging




Figure 3: Cost vs accuracy of the three monitoring policies. The DCI selector retains 99% of the multi-D-over-1-D AUC gain at 38% of the multi-D per-window cost.

over [1.22, 2.41]: real drift is predominantly high-complexity (90% of blocks above the frontier), with a 10% minority of low-complexity blocks. The heterogeneous regime where the multi-dimensional detector is required is thus not a synthetic artifact but the common case on this trace — and the honest reading cuts both ways. Routed at $\tau = 1.5$, the selector prices the trace at 2.12 ms per window, 90% of the always-multi-D monitoring cost: on a workload like this the probe's value is not a large saving but *foreknowledge* — it tells the operator, before any detector runs, that this workload must pay for full resolution, and exactly which tenth of it need not. (The trace carries no ground-truth onsets, so detection quality is not scored here; costs use the per-window timings of §6.)

6.3 The selector and the monitoring cost (RQ3)

The multi-dimensional detector costs 2.35 ms per window against the 1-D detector's 0.0087 ms — a cost ratio of $268.7 \times \pm 5.4$ (mean $\pm$ SD of the per-repeat ratios over the five timing repeats). Since one detector runs on *every* window, this is the always-on monitoring overhead. Figure 3 places the three policies in cost-accuracy space. The always-multi-D detector reaches mean AUC 0.970 at 2.35 ms; always-1-D reaches 0.910 at 0.0087 ms; the DCI selector, routing 62% of the 450 runs to the 1-D detector (six of the nine cells by mean DCI; runs of a low-DCI cell whose block DCI crosses τ fall back to the multi-dimensional detector), reaches 0.970 at 0.89 ms. The selector thus retains 99% (bootstrap 95% CI 90–104%; $B=10^4$, stratified by cell) of the multi-D-over-1-D detection-AUC gain at 38% (95% CI 36–40%) of the multi-D monitoring cost — a retained-gain interval that spans 100%, so the selector is statistically indistinguishable from capturing the entire gain.

The saving at scale. These costs are *per window*, and the monitoring layer evaluates one detector on *every* window, so at a window-arrival rate of R windows/s it consumes Rc of CPU — the routing choice scales with load. Taking $R = 100$ windows/s as a reference, the always-multi-D layer occupies 23.5% of a core, always-1-D 0.09%, and the DCI selector 8.9%; DCI-routing thus frees $\approx 14.6\%$ of a core relative to the full kernel *at the full kernel's detection quality*. On a long-lived, high-rate workload it is this always-on overhead — not the one-off cost of an advisor call — that the monitoring layer minimises; where windows are infrequent the choice is moot, the honest boundary of the result.




Figure 4: Threshold sensitivity. Retained multi-D-over-1-D AUC-gain (left axis) and monitoring cost as a fraction of the multi-dimensional detector (right axis) versus the routing threshold τ , over all 450 runs. The shaded band is the empty per-cell DCI interval (1.41, 1.75); $\tau \approx 1.5$ sits at the knee where the selector keeps $\geq 99\%$ of the gain at $\approx 38\%$ of the cost.




Figure 5: Why $\tau \approx 1.5$ is principled: detection AUC of the multi-dimensional detector and the best single axis versus DCI, in a synthetic model matched to the two detectors (rank-one drift along the hardest, maximally spread direction; iid steady noise). The curves separate at DCI ≈ 1.5 : the deployed threshold sits at the onset of the detectability gap.

6.4 Threshold sensitivity

The routing frontier $\tau \approx 1.5$ is not a tuned constant. Figure 4 sweeps it over all 450 runs. As τ falls toward 1 every window routes to the multi-dimensional detector — full accuracy at full cost; as τ rises past the high-complexity cells the mixed-drift gain is lost. Between the two lies a plateau: for $\tau \in [1.2, 1.5]$ the selector retains $\geq 99\%$ of the multi-D-over-1-D gain while paying only 37–48% of the multi-D cost, and the retained gain collapses only once τ crosses 1.75 and the *mixed* cells are misrouted (63.9% retained at $\tau=1.75$, 21.1% at $\tau=2.0$). At the cell level no cell's mean DCI lies inside the open band (1.41, 1.75) — its endpoints are themselves the two nearest cell means — so any τ within it routes all nine cells identically. The frontier is therefore robust — $\tau \approx 1.5$ sits at the efficient knee, the largest (cheapest) threshold that still captures essentially all the gain.

Why the frontier is principled. The location of τ is not an accident of these workloads. In a synthetic model matched to the

two detectors—drift of graded magnitude along the *hardest*, maximally spread direction, under iid and correlated steady noise—the multi-D-over-1-D detectability gap tracks DCI: at most 0.025 AUC below $\text{DCI} \approx 1.25$, climbing steeply across the empty band (0.04–0.06 over $\text{DCI} 1.3\text{--}1.45$), and largest (≈ 0.07) over $\text{DCI} \approx 1.5\text{--}1.9$, precisely the band the router escalates (Figure 5). The threshold thus sits on the steep rise of the detectability gap, just below its peak—model evidence that $\tau \approx 1.5$ marks where multi-dimensional fusion pays most, not a constant tuned to these cells. (At the sweep’s far end, where ever-weaker drift inflates DCI toward the noise floor, both detectors degrade together: that region is noise-dominated, unlike the strong heterogeneous drift of the mixed cells, whose multi-dimensional AUC stays 0.91–1.00.) The same sweep probed the converse risk, a low-DCI drift hidden *off-axis*, and found none: a drift with DCI low enough to route to the 1-D detector must dominate the steady noise, and a drift that strong remains detectable on a single axis even when maximally spread. Low DCI genuinely implies 1-D detectability in this model; the frontier’s residual dependence on the steady-noise structure is discussed in §8.

6.5 The tolerance, and further checks

Three checks support the result. (i) The composite is not a detector: the scalar HSM’s detection AUC falls to 0.68 where a single axis is perfect, confirming a weighted average dilutes single-axis drift. (ii) The tolerance δ_d is verified, not assumed: against known-truth synthetic data its DeLong standard error matches a stratified bootstrap to within 6% and its 95% interval covers at 0.94–0.96 (inverse-variance pooling, biased for AUC, covers only 51%). (iii) The regime map is recovered from as few as 10 seeds, so 50 is ample.

6.6 Ablation: why the participation ratio?

DCI is one of several ways to summarise a spectrum’s concentration; does the routing verdict depend on the choice? Table 3 recomputes, per cell, three alternatives: the *effective rank* $\exp(-\sum_i p_i \ln p_i)$ (the Hill number of order one [11]); the *stable rank* $\text{tr}(C)/\lambda_{\max}$; and the dominant-mode share $p_1 = \lambda_{\max}/\text{tr}(C)$. Every one separates the low-complexity (template/volume) cells from the high (mixed) cells with an empty band (for DCI, on (1.41, 1.75); likewise for the others), so the verdict is robust to the choice of functional. The participation ratio is nonetheless the right runtime statistic: it is the *only* one of the four computable without an eigendecomposition — a trace and a Frobenius norm ($O(D^2)$; §4), whereas effective rank needs the whole spectrum and stable rank and p_1 need the top eigenvalue ($O(D^3)$ on every routed block). Same verdict, lowest cost.

6.7 The whitened refinement

DCI reads the raw drift covariance; the steady noise Σ is left to the detectors. The noise-aware refinement — DCI_w , the participation ratio of $\Sigma^{-1/2} C \Sigma^{-1/2}$ — describes the drift a *whitened* detector family sees. On the identical trajectories and seeds as the regime map (the raw column reproduces Table 2 to the printed digit), whitening only *concentrates* the drift: $\text{DCI}_w \leq \text{DCI}$ on all nine cells, collapsing the volume cells to ≈ 1.01 and JOB-mixed from 1.92 to 1.08, while pgbench-mixed is unchanged at 1.75 — its Σ is near-isotropic, so whitening is a no-op there. The statistic is stable over the covariance ridge in $[10^{-8}, 10^{-6}]$; larger ridges interpolate back toward the raw value, as un-whitening should.

The detector this concentrated spectrum promises exists: the whitened *matched filter* — the squared projection of the whitened deviation on the dominant whitened drift mode, the strictly-more-powerful 1-D detector the white-noise model of Theorem 1 anticipates. Under the map’s own tolerance rule, it is sufficient in *every* cell — including the three mixed cells, where the best raw axis loses 0.16–0.21 AUC yet the whitened filter comes within 0.015, 0.001, and 0.035 of the multi-dimensional detector (JOB, TPC-H, and pgbench; the last with the loosest floor, on the noisiest cell). This does not create a third routing option, because the whitened filter is the full-kernel tier in disguise: $z = \Sigma^{-1/2}(f - \mu)$ needs all five axes of f on every window — the 2.35 ms feature cost — whereas the cheap tier’s 0.0087 ms comes precisely from computing a *single* axis. The monitoring layer thus has two cost tiers, not three: axis-restricted (cheap) and full-kernel (expensive, in either its Mahalanobis or matched-filter form). Raw DCI routes *between* the tiers — erring only toward escalation, since $\text{DCI}_w \leq \text{DCI}$ — while DCI_w describes structure *within* the full tier. Table 2 is therefore the map for the detector family a cheap tier can actually deploy.

6.8 The regime is not tied to the kernel

DCI reads only the covariance of a per-window feature vector (§4), so its regime should not depend on the particular similarity axes it is fed. We test this directly: on the *same* drift windows and seeds, we recompute DCI and the 1-D/multi-D detection AUC under three representations that use *none* of the HSM kernel’s machinery (Table 4)—a generic three-axis similarity (adjacent-window template cosine, query-count ratio, table Jaccard), a raw per-window template-frequency vector, and a table-incidence bag.

Two findings. First, the generic three-axis similarity—which shares no construction with the HSM-style feature—reproduces the regime almost exactly (single-cause drift low-complexity, mixed high-complexity, and the same frontier $\tau \approx 1.5$): the low-complexity of ordinary drift is a property of the drift, not of that kernel. Second, the routing rule is valid in *every* representation. Wherever DCI is low the 1-D detector matches the multi-dimensional one; wherever DCI is high the 1-D detector falls and the multi-dimensional detector recovers it—including the raw representations, where even ordinary *template* drift turns high-complexity (DCI up to 4.0) because it moves many raw dimensions, and DCI correctly routes it to the multi-dimensional detector. What changes across representations is therefore not the router’s validity but the *operating point*: a bounded-similarity featurization places ordinary drift at low complexity, so the cheap detector usually suffices; a raw high-dimensional one does not. DCI predicts detector sufficiency on whatever features the monitoring layer computes.

6.9 End-to-end deployment (RQ4)

To confirm the routed detector drives a real adaptation loop, we connect the DCI-gated selector to a live, *off-the-shelf* index advisor—Dexter 0.6.3, a HypoPG-based recommender (HypoPG 1.4.2), on PostgreSQL 16.9 (TPC-H, SF 1), and an equality-sort-range recommender on MongoDB 8.0.8 (a 5M-document workload)—over 10 randomized complete-block (RCB) replicates per drift schedule, with seeded schedules whose onsets are ground truth. The advisor is a black-box downstream consumer: the selector decides only *when* to invoke it, so any advisor serves. We therefore score the gate’s own firings against the ground

Table 3: Ablation: the participation ratio (DCI) against three alternative spectrum-concentration measures, per cell (mean over 50 seeds, all three workloads). Every measure separates the LOW (template/volume) cells from the HIGH (mixed) cells; only DCI needs no eigendecomposition. All functionals are computed per seed and then averaged, so the tabulated stable rank need not equal the reciprocal of the tabulated p_1 .

Workload	Drift	DCI	eff. rank	stable rank	p_1
TPC-H	template	1.02	1.07	1.01	0.99
TPC-H	volume	1.41	1.57	1.22	0.83
TPC-H	mixed	1.97	2.17	1.62	0.63
JOB	template	1.11	1.22	1.05	0.95
JOB	volume	1.13	1.26	1.06	0.94
JOB	mixed	1.92	2.11	1.62	0.63
pgbench	template	1.03	1.08	1.02	0.98
pgbench	volume	1.01	1.02	1.00	1.00
pgbench	mixed	1.75	1.85	1.52	0.68

Table 4: DCI and 1-D/multi-D detection AUC under three *kernel-free* feature representations (pooled over the three SQL workloads, 50 seeds; each cell reads “DCI 1D/mD”). A bounded-similarity representation reproduces the regime of §6.2; raw representations move ordinary drift to high complexity—but in all of them high DCI coincides with 1-D insufficiency, so the DCI router stays valid.

representation	template	volume	mixed
generic 3-axis sim.	1.11 .93/.93	1.04 1.0/1.0	1.89 .79/.97
raw template-freq.	4.02 .78/.95	1.87 .97/1.0	3.18 .82/.97
table-incidence bag	2.25 .80/.86	1.03 .97/.97	1.81 .87/.91

truth—detection recall, invocations, false alarms—not the advisor’s economic value (§7).

On both engines the routed advisor fires on *every* onset at zero lag (recall 1.00; Table 6), so the detection $\rightarrow$ advisor $\rightarrow$ index chain closes end-to-end, at query latency comparable to an always-1-D gate (latency is dominated by execution, not the sub-millisecond detector). This is the low-complexity regime of §6.2 seen live: ordinary drift is *near* rank-one, and the cheap 1-D detector catches it on both a relational and a document store. The DCI-gated selector moreover fires *fewer* false alarms than always-1-D — 3.1 vs 3.9 on MongoDB (paired Wilcoxon $p=0.016$; 0.6 vs 1.2 on PostgreSQL, same direction) — by escalating to the more precise multi-dimensional detector on the windows DCI flags. Replaying the logged feature stream through the frozen gate, 14.6% of PostgreSQL and 20.8% of MongoDB deployment windows ran the multi-dimensional detector: the first two windows of each block default to it while the drift covariance is not yet estimable, and beyond that warm-up 6.8% and 13.6% were DCI-triggered escalations — the gate stays on the cheap detector for the large majority of windows and buys its precision exactly where DCI asks for it. where that detector is *required* (the mixed cells, §6.2) is not exercised by these ordinary workloads; we exercise it next — first offline at the firing level, then live on both engines.

The heterogeneous regime at the firing level. The mixed cells can be exercised at the gate’s own granularity offline: replaying the same drift trajectories (50 seeds per cell, all three workloads) through three gate configurations — always-multi-D ($\tau=0$), the DCI-gated selector ($\tau=1.5$), and always-1-D ($\tau \rightarrow \infty$) — under strict onset scoring (a hit only if the gate fires exactly at the transition window; every other fire is a false alarm). On *mixed* drift the always-1-D gate reaches recall 0.69 against the multi-D gate’s 0.93 — the guaranteed insufficiency of Corollary 2, now visible in

firing decisions — while the DCI-gated selector recovers 0.90 at 45% of the multi-D monitoring cost, with matched false alarms (7.1 vs 7.2 per run). On single-cause drift all three configurations tie (recall 0.89–1.00) and the selector pays 6–42% of the multi-D cost.

The heterogeneous regime live. We then run the same three configurations against the live advisors on a mixed-drift schedule that reproduces the mixed cells of §6.2 in production form: six phases of four windows per 24-window block, onsets alternating template/volume — a template move half-swaps the active template set (PostgreSQL) or switches the dominant document modality (MongoDB); a volume move toggles the per-window query count (20 $\leftrightarrow$ 32 relational, 20 $\leftrightarrow$ 48 document) with the mix held. These blocks sit where the regime map puts mixed drift: pooled onset DCI 2.04 (PostgreSQL) and 1.72 (MongoDB) in offline pre-validation, and the live gate agrees — its rolling DCI runs below τ until the first onset (medians 1.21/1.15) and at 1.7–2.4 from the first volume move on.

The sweep instantiates the domain boundary of §7 in production, identically on both engines (Table 5). Below the boundary the cheap detector is the right tool: every configuration fires on the opening template onset (window 5, rolling DCI ≈ 1.2). Beyond it, the frozen 1-D detector locks onto the direction that dominates its whitened trajectory — here the volume axis — and from that point fires on every volume onset (20/20 per engine) while missing every later template onset (0/20 per engine, never recovered a window late): recall 0.60 on both stores, the insufficiency of Corollary 2 at the decision level, with the same deterministic signature (hits at windows 5, 9, 17; misses at 13, 21; in 10/10 blocks) on a relational and a document engine. The DCI-gated selector escalates across the boundary and matches the always-multi-D reference decision for decision (99.2%/98.3% of windows; recall 1.00 on both onset types) at 82%/81% of its monitoring cost — on a stream built so the expensive detector is needed most of the time; single-cause drift pays 6–42% (above).

Whether a missed onset costs query latency depends on who owns the indexes. On PostgreSQL the structural backbone already covers the mixed queries’ access paths and the damage is nil (post-onset/steady latency contrast 1.00) — an honest null. On MongoDB the advisor owns each modality’s index: after the missed template onsets the document phase runs unindexed at 10.9 $\times$ the reference latency (geometric mean over the 30 stale windows) until the next volume onset — one the 1-D detector *can* see — repairs the indexes as a side effect; the pre-registered contrast summarizes this conservatively at 1.23, since damage persists past its +2-window horizon. The gated selector’s contrast

Table 5: Live mixed drift: three detector resolutions in front of the same advisors, 10 RCB blocks per engine, onsets alternating template (T, 30 total) and volume (V, 20 total); the hit/miss pattern is deterministic across blocks. FA = fires per block outside an onset or its +1 window; cost = per-window monitoring cost as % of always-multi-D; fire=m-D = share of windows whose firing decision equals always-multi-D's; lat. = post-onset/steady latency contrast vs. always-multi-D (the MongoDB always-1-D stale stretch runs 10.9× before its accidental repair).

Engine	Gate	T	V	FA	Cost	fire=m-D	lat.
PostgreSQL	always-multi-D	30	20	0.2	100%	—	—
PostgreSQL	DCI-gated	30	20	0.3	82%	99.2%	1.01
PostgreSQL	always-1-D	10	20	0.6	8.7%	87.9%	1.00
MongoDB	always-multi-D	30	20	0.8	100%	—	—
MongoDB	DCI-gated	30	20	1.0	81%	98.3%	1.00
MongoDB	always-1-D	10	20	1.3	8.7%	83.8%	1.23

Table 6: End-to-end: the DCI-gated advisor vs an always-1-D gate on live PostgreSQL and MongoDB, 10 RCB blocks per row, scored against the seeded ground-truth drift ($\pm 95\%$ CI). Both fire on every onset at zero lag; DCI-gating issues fewer false alarms.

Engine	Gate	Recall	Advisor calls	False alarms
PostgreSQL	DCI-gated	1.00	4.1±0.6	0.6±0.3
PostgreSQL	always-1-D	1.00	4.7±0.8	1.2±0.5
MongoDB	DCI-gated	1.00	7.5±1.5	3.1±0.9
MongoDB	always-1-D	1.00	8.3±1.5	3.9±1.0

is 1.00 on both engines: it built the same indexes as the reference. (Raw cross-configuration latency ratios are not read as effects: the τ legs run sequentially, so machine-state drift inflates a leg uniformly; the within-leg contrast cancels it, and the gated leg is its negative control — raw 1.06 on PostgreSQL, contrast 1.01.)

7 Discussion and Future Work

What the result says. The empirical core — ordinary workload drift is *low-complexity*, occupying essentially a single mode, while only heterogeneous drift spreads across several — is a statement about how workloads move, not merely about detectors: it is why a nearly free always-on layer suffices almost always, and the bounds of §4 say when it will not. Reading only the covariance of a per-window feature vector, DCI is indifferent to *how* that vector is produced and *what* its verdict triggers — a general monitoring-resolution knob for self-driving data systems, of which index-advisor gating is one instance.

The domain boundary, drawn by the primitive itself. The scope of the cheap path is not asserted; it is measured and proved. *Covered:* abrupt-onset, single-cause drift—the low-DCI zone where Corollary 1 guarantees the 1-D detector retains a $1 - \delta$ share of the full power—demonstrated across three relational workloads and a document store, and live on two engines (§6.9). *Covered by escalation:* heterogeneous drift. Corollary 2 proves no single direction suffices once $DCI > (1 - \delta)^{-2}$; the gate detects exactly this condition and escalates to the multi-dimensional detector, so the system does not fail there—§6.9 shows this live on both engines (recall 1.00 across the boundary against the fixed detector's 0.60)—what this paper does not develop is that path's end-to-end economics. *Outside the domain:* gradual (non-abrupt) drift, for

which the position detectors were not designed (§5); arrival-only drift, which this kernel's S_P cannot isolate (§6); and correlated-noise refinements of the frontier itself (§8). A system built on the layer inherits this boundary—and can read it, in advance, off the same number.

Future work. Four directions follow. (i) **Economics of the multi-dimensional path.** §6.9 validates the routed advisor live in both regimes: the cheap detector where Corollary 1 covers it, escalation where Corollary 2 proves a fixed cheap detector cannot go. What remains open is not *whether* the multi-dimensional path is needed, nor whether the gate reaches it (it does, decision-matched to always-multi-D), but its end-to-end *economics* — plan quality, invocations avoided, and when paying its cost is worthwhile — the journal-scale sequel. (ii) **Representations.** Re-derive the regime map under alternative per-window features (a raw query-frequency vector, or learned query/plan embeddings) to test how far the frontier $\tau \approx 1.5$ travels and to place DCI on modern learned workload models. (iii) **Engines and data models.** The regime map already transfers to a document store (§6.2); extending it to column and vector stores, and to heterogeneous workload *types* beyond the controlled drift model, is the natural next step. (iv) **Adaptivity and composition.** Let τ self-tune online as a workload's drift geometry shifts, and compose DCI's *select-on-complexity* routing with the *escalate-on-event* hierarchical testing of §2 [36, 37] into a single two-layer monitor.

8 Threats to Validity

Scope: a compact end-to-end. The live deployment (§6.9) is a 10-block-per-regime validation on two engines — detection recall, advisor invocations, and query latency — not a full economic study of plan quality and invocations avoided across workloads; that study is future work (§7). Because the three detector configurations run sequentially, raw cross-configuration latency ratios include machine-state drift; the mixed-regime deployment therefore reads latency through the within-configuration post-onset/steady contrast, whose negative control (the gated leg, which builds the reference's indexes) comes out at 1.01.

Online deployment. The end-to-end validation runs the gate over controlled blocks with a fixed threshold τ against a seeded ground-truth schedule. A fully *online* deployment—the gate running continuously on a production stream, τ self-tuning as the workload's drift geometry shifts, and the gate robust to drift in its own steady-state statistics—is future work (§7), as is the advisor's economic payoff under live gating.

Construct: synthetic drift. Drift trajectories come from a generator anchored on the phase profiles of a real workload trace; they span template, volume, and mixed drift but are not an in-the-wild deployment. We corroborate the analytic workloads with an OLTP one (pgbench) and a full document-store regime map on MongoDB (§6.2), but a field study remains future work. The model is also abrupt-onset by construction—phase transitions produce the transient dips the position detectors are built for (§5); gradual or continuous drift, where motion-based statistics become competitive, lies outside the validated domain (§7).

External: the frontier and the feature set. The regime map is demonstrated on one workload-similarity kernel. DCI depends only on the feature covariance (§3), so the construction transfers, but the frontier $\tau \approx 1.5$ is empirical across three SQL workloads and a document store, not a proven constant — a different feature

set could place it elsewhere (§7). DCI also reads only the drift covariance, not the steady-state *noise* covariance Σ ; the frontier is therefore Σ -dependent, and a noise-aware refinement — the participation ratio of the *whitened* drift covariance $\Sigma^{-1/2}C\Sigma^{-1/2}$ — is a natural next step. That refinement is measured in §6.7: $\text{DCI}_w \leq \text{DCI}$ on every cell, so the deployed raw-DCI router errs only toward escalation, and the refinement’s own matched filter prices as the full-kernel tier — the frontier’s Σ -dependence bounds *which* cheap axis suffices, not whether the two-tier routing stands.

Internal: cost measurement. Absolute per-window costs are single-machine wall-clock on one Apple M4 platform and will shift with hardware and library versions; the 268.7 $\times$ cost ratio and the DCI/regime-map results are platform-independent by construction, since DCI is dimensionless and the map is defined on AUC and DCI alone.

9 Conclusion

Workload drift has a measurable complexity — the effective number of independent modes its motion occupies — and the Drift Complexity Index estimates it in $O(D^2)$ with no eigendecomposition. We proved that DCI bounds detector sufficiency — the best 1-D detector’s power is the DCI-bracketed dominant-mode share — and confirmed it empirically: ordinary drift is *near* rank-one, so one cheap dimension suffices, while only heterogeneous drift needs the full detector. That regime is reproduced by a kernel-free similarity and exhibited by a real query-log trace, so it is a property of the drift, not of one kernel. Reading this off *before* any detector runs, DCI routes the always-on monitoring layer to the cheapest sufficient detector, retaining 99% of the full kernel’s detection-AUC gain at 38% of its per-window cost. The practitioner rule is short: measure a workload’s DCI, then deploy the detector its regime — and the bound — call for. DCI is a monitoring-resolution primitive for self-driving data systems; advisor gating is its first application.

10 Artifacts

All code, data, and run artifacts reproducing every number, figure, and table in this paper are publicly available at <https://github.com/stevezhai139/dci-reproducibility>. The repository contains the workload-similarity kernel and the DCI implementation, the drift generators and regime-map pipeline (§6), the kernel-free representations and the SkyServer adapter (§6.8), the DCI-routed gate and both live end-to-end harnesses with their ordinary and mixed drift schedules (§6.9), and the official run outputs with per-run provenance (seeds, environment lock, engine and driver versions). REPRODUCE.md maps each paper artifact to its script, its official run directory, and the expected output; a read-only preflight script verifies the live environment (PostgreSQL 16.9 with HypoPG and Dexter; MongoDB 8.0.8) before any database experiment. The offline pipeline is fully deterministic — every cell is keyed by a stable per-seed hash — and requires no database.

Acknowledgments

Portions of this paper’s code and text were developed with the assistance of a generative AI system (Anthropic’s Claude): the experiment-harness variants, run and analysis scripts, and initial drafts of parts of the text were AI-generated under the authors’ direction and subsequently reviewed, edited, and validated by the authors. All experimental results were measured by the authors

on the pinned environment of §6 from the released repository; references were selected and verified by the authors.

References

- [1] Sanjay Agrawal, Surajit Chaudhuri, Lubor Kolár, Arnd Christian Marathe, Vivek R. Narasayya, and Manoj Syamala. 2004. Database Tuning Advisor for Microsoft SQL Server 2005. In *Proc. 30th International Conference on Very Large Data Bases (VLDB)*. Morgan Kaufmann, Toronto, Canada, 1110–1121. doi:10.1016/B978-012088469-8.50097-8
- [2] Shruti Arora et al. 2024. A Systematic Review on Detection and Adaptation of Concept Drift in Streaming Data Using Machine Learning Techniques. *WIRES Data Mining and Knowledge Discovery* 14, 3 (2024), e1536. doi:10.1002/widm.1536
- [3] Albert Bifet and Ricard Gavaldà. 2007. Learning from Time-Changing Data with Adaptive Windowing. In *Proc. SIAM International Conference on Data Mining (SDM)*. 443–448. doi:10.1137/1.9781611972771.42
- [4] Matteo Brucato, Tarique Siddiqui, Wentao Wu, Vivek Narasayya, and Surajit Chaudhuri. 2024. Wred: Workload Reduction for Scalable Index Tuning. *Proc. ACM on Management of Data (SIGMOD)* 2, 1 (2024).
- [5] Surajit Chaudhuri and Vivek R. Narasayya. 1997. An Efficient, Cost-Driven Index Selection Tool for Microsoft SQL Server. In *Proc. 23rd International Conference on Very Large Data Bases (VLDB)*. 146–155.
- [6] Surajit Chaudhuri and Vivek R. Narasayya. 1998. AutoAdmin “What-If” Index Analysis Utility. In *Proc. ACM SIGMOD International Conference on Management of Data*. 367–378. doi:10.1145/276304.276337
- [7] Elizabeth R. DeLong, David M. DeLong, and Daniel L. Clarke-Pearson. 1988. Comparing the Areas under Two or More Correlated Receiver Operating Characteristic Curves: A Nonparametric Approach. *Biometrics* 44, 3 (1988), 837–845. doi:10.2307/2531595
- [8] João Gama, Pedro Medas, Gladys Castillo, and Pedro Pereira Rodrigues. 2004. Learning with Drift Detection. In *Advances in Artificial Intelligence – SBIA 2004 (Lecture Notes in Artificial Intelligence, Vol. 3171)*. 286–295. doi:10.1007/978-3-540-28645-5_29
- [9] João Gama, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. 2014. A Survey on Concept Drift Adaptation. *Comput. Surveys* 46, 4, Article 44 (2014). doi:10.1145/2523813
- [10] Peiran Gao, Eric Trautmann, Byron Yu, Gopal Santhanam, Stephen Ryu, Krishna Shenoy, and Surya Ganguli. 2017. A Theory of Multineuronal Dimensionality, Dynamics and Measurement. doi:10.1101/214262 bioRxiv 214262.
- [11] M. O. Hill. 1973. Diversity and Evenness: A Unifying Notation and Its Consequences. *Ecology* 54, 2 (1973), 427–432. doi:10.2307/1934352
- [12] Hanxian Huang, Tarique Siddiqui, Rana Alotaibi, Carlo Curino, Jyoti Leeka, Alekh Jindal, Jesús Camacho-Rodríguez, Yuanyuan Tian, and Yi Zhao. 2024. Sibyl: Forecasting Time-Evolving Query Workloads. *Proc. ACM on Management of Data* 2, 1, Article 41 (2024). doi:10.1145/3639308
- [13] Tze Leung Lai. 1995. Sequential Change-point Detection in Quality Control and Dynamical Systems. *Journal of the Royal Statistical Society: Series B* 57, 4 (1995), 613–658. doi:10.1111/j.2517-6161.1995.tb02052.x
- [14] Tze Leung Lai. 1998. Information Bounds and Quick Detection of Parameter Changes in Stochastic Systems. *IEEE Transactions on Information Theory* 44, 7 (1998), 2917–2929. doi:10.1109/18.737522
- [15] Jiale Lao, Yibo Wang, Yufei Li, Jianping Wang, Yunjia Zhang, Zhiyuan Cheng, Wanghu Chen, Mingjie Tang, and Jianguo Wang. 2024. GPTuner: A Manual-Reading Database Tuning System via GPT-Guided Bayesian Optimization. *Proc. VLDB Endowment* 17, 8 (2024), 1939–1952.
- [16] Patrick A. Lee and T. V. Ramakrishnan. 1985. Disordered Electronic Systems. *Reviews of Modern Physics* 57, 2 (1985), 287–337. doi:10.1103/RevModPhys.57.287
- [17] Guanli Liu and Renata Borovica-Gajic. 2025. DriftBench: Defining and Generating Data and Query Workload Drift for Benchmarking. arXiv:2510.10858 [cs.DB]
- [18] Jie Lu, Anjin Liu, Fan Dong, Feng Gu, João Gama, and Guangquan Zhang. 2019. Learning under Concept Drift: A Review. *IEEE Transactions on Knowledge and Data Engineering* 31, 12 (2019), 2346–2363. doi:10.1109/TKDE.2018.2876857
- [19] Lin Ma, Dana Van Aken, Ahmed Hefny, Gustavo Mezerhane, Andrew Pavlo, and Geoffrey J. Gordon. 2018. Query-based Workload Forecasting for Self-Driving Database Management Systems. In *Proc. ACM SIGMOD International Conference on Management of Data*. 631–645. doi:10.1145/3183713.3196908
- [20] Michael J. Mior, Kenneth Salem, Ashraf Aboulnaga, and Rui Liu. 2016. NoSE: Schema Design for NoSQL Applications. In *Proc. IEEE 32nd International Conference on Data Engineering (ICDE)*. 181–192. doi:10.1109/ICDE.2016.7498240
- [21] Jerzy Neyman and Egon S. Pearson. 1933. On the Problem of the Most Efficient Tests of Statistical Hypotheses. *Philosophical Transactions of the Royal Society A* 231 (1933), 289–337. doi:10.1098/rsta.1933.0009
- [22] Rafiullah Omar, Justus Bogner, Joran Leest, Vincenzo Stoico, Patricia Lago, and Henry Muccini. 2024. How to Sustainably Monitor ML-Enabled Systems? Accuracy and Energy Efficiency Tradeoffs in Concept Drift Detection. In *Proc. 10th International Conference on ICT for Sustainability (ICT4S)*. doi:10.1109/ICT4S64576.2024.00010
- [23] E. S. Page. 1954. Continuous Inspection Schemes. *Biometrika* 41, 1/2 (1954), 100–115. doi:10.1093/biomet/41.1-2.100

- [24] Zhicheng Pan et al. 2025. Hyper: Hybrid Physical Design Advisor with Multi-agent Reinforcement Learning. In *Proc. IEEE 41st International Conference on Data Engineering (ICDE)*. 1565–1578.
- [25] Andrew Pavlo, Gustavo Angulo, Joy Arulraj, Haibin Lin, Jiexi Lin, Lin Ma, Prashanth Menon, Todd C. Mowry, Matthew Perron, Ian Quah, Siddharth Santurkar, Anthony Tomasic, Skye Toor, Dana Van Aken, Ziqi Wang, Yingjun Wu, Ran Xian, and Tieying Zhang. 2017. Self-Driving Database Management Systems. In *Proc. 8th Biennial Conference on Innovative Data Systems Research (CIDR)*.
- [26] R. Malinga Perera, Bastian Oetomo, Benjamin I. P. Rubinstein, and Renata Borovica-Gajic. 2021. DBA Bandits: Self-Driving Index Tuning under Ad-hoc, Analytical Workloads with Safety Guarantees. In *Proc. IEEE 37th International Conference on Data Engineering (ICDE)*. 600–611. doi:10.1109/ICDE51399.2021.00058
- [27] R. Malinga Perera, Bastian Oetomo, Benjamin I. P. Rubinstein, and Renata Borovica-Gajic. 2023. No DBA? No Regret! Multi-Armed Bandits for Index Tuning of Analytical and HTAP Workloads with Provable Guarantees. *IEEE Transactions on Knowledge and Data Engineering* 35, 12 (2023), 12855–12872. doi:10.1109/TKDE.2023.3271664
- [28] Alfréd Rényi. 1961. On Measures of Entropy and Information. In *Proc. 4th Berkeley Symposium on Mathematical Statistics and Probability*, Vol. 1. 547–561.
- [29] A. Reungsilpulkarn, A. Sangmanee, and C. Phiphatkul. 2026. Adaptive Indexing for Schema-Less Temporal Databases Using Hierarchical Similarity Measurement. In *Proc. 9th Int. Conf. on Information and Computer Technologies (ICICT)*. Honolulu, HI, USA. doi:10.1145/3803291.3803321
- [30] A. Reungsinkonkarn, N. Kasamsuksomboon, and S. Rasniyom. 2026. HSM-Based Workload Similarity Analysis for Temporal Database Indexing. In *Proc. 2026 International Electrical Engineering Congress (iEECON)*. Pattaya, Chonburi, Thailand.
- [31] Dana Van Aken, Andrew Pavlo, Geoffrey J. Gordon, and Bohan Zhang. 2017. Automatic Database Management System Tuning Through Large-scale Machine Learning. In *Proc. ACM SIGMOD International Conference on Management of Data*. 1009–1024. doi:10.1145/3035918.3064029
- [32] Junxiong Wang, Immanuel Trummer, and Debabrota Basu. 2021. UDO: Universal Database Optimization using Reinforcement Learning. *Proc. VLDB Endowment* 14, 13 (2021), 3402–3414. doi:10.14778/3484224.3484236
- [33] Xiaoying Wang, Wentao Wu, Vivek R. Narasayya, and Surajit Chaudhuri. 2025. Esc: An Early-Stopping Checker for Budget-aware Index Tuning. *Proc. VLDB Endowment* 18, 5 (2025), 1278–1290.
- [34] Zijia Wang, Haoran Liu, Chen Lin, Zhifeng Bao, Guoliang Li, and Tianzheng Wang. 2024. Leveraging Dynamic and Heterogeneous Workload Knowledge to Boost the Performance of Index Advisors. *Proc. VLDB Endowment* 17, 7 (2024), 1642–1654. doi:10.14778/3654621.3654631
- [35] Chenning Wu, Sifan Chen, Wentao Wu, Yinan Jing, Zhenying He, Kai Zhang, and X. Sean Wang. 2026. UTune: Towards Uncertainty-Aware Online Index Tuning. arXiv:2601.18199.
- [36] Shujian Yu and Zubin Abraham. 2017. Concept Drift Detection with Hierarchical Hypothesis Testing. In *Proc. SIAM International Conference on Data Mining (SDM)*. 768–776. doi:10.1137/1.9781611974973.86
- [37] Shujian Yu, Xiaoyang Wang, and José C. Principe. 2018. Request-and-Reverify: Hierarchical Hypothesis Testing for Concept Drift Detection with Expensive Labels. In *Proc. 27th International Joint Conference on Artificial Intelligence (IJCAI)*. 3033–3039. doi:10.24963/ijcai.2018/421
- [38] Wei Zhou, Chen Lin, Xuanhe Zhou, and Guoliang Li. 2024. Breaking It Down: An In-Depth Study of Index Advisors. *Proc. VLDB Endowment* 17, 10 (2024), 2405–2418. doi:10.14778/3675034.3675035